

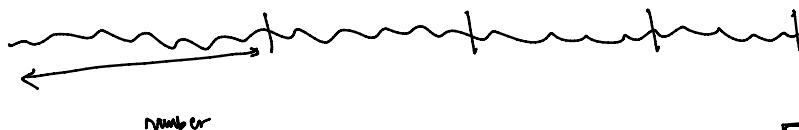
Boosting

21 February 2022 20:45

16

Clustering

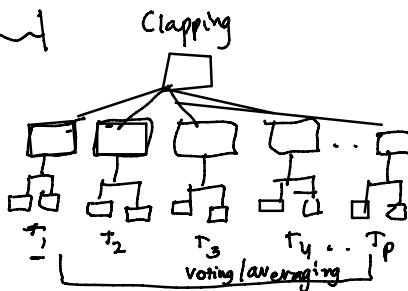
(avg, min, max)



Bagging - Parallel Process

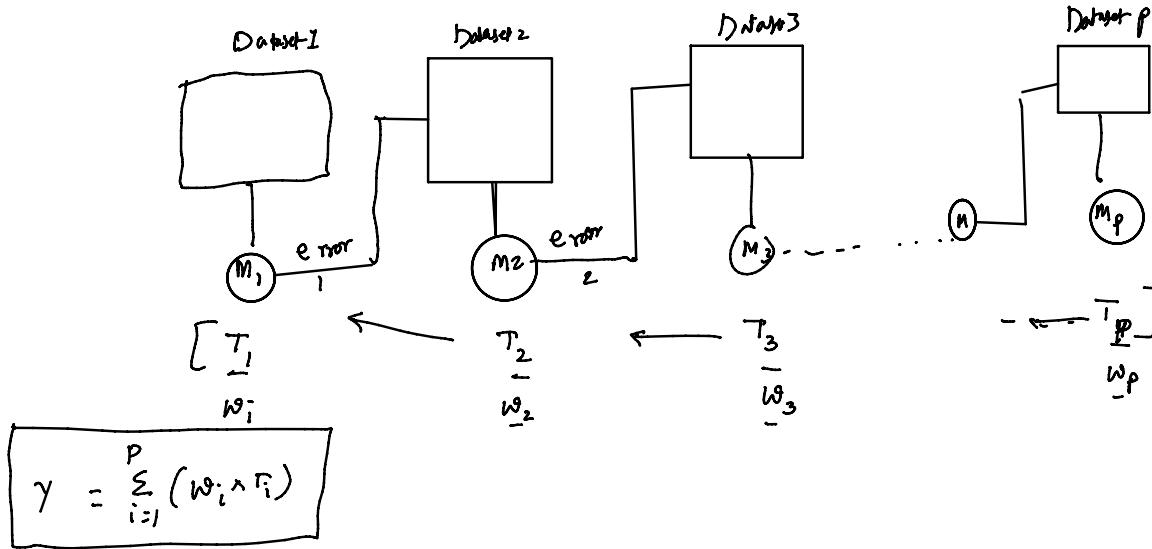
label

Clapping



Boosting

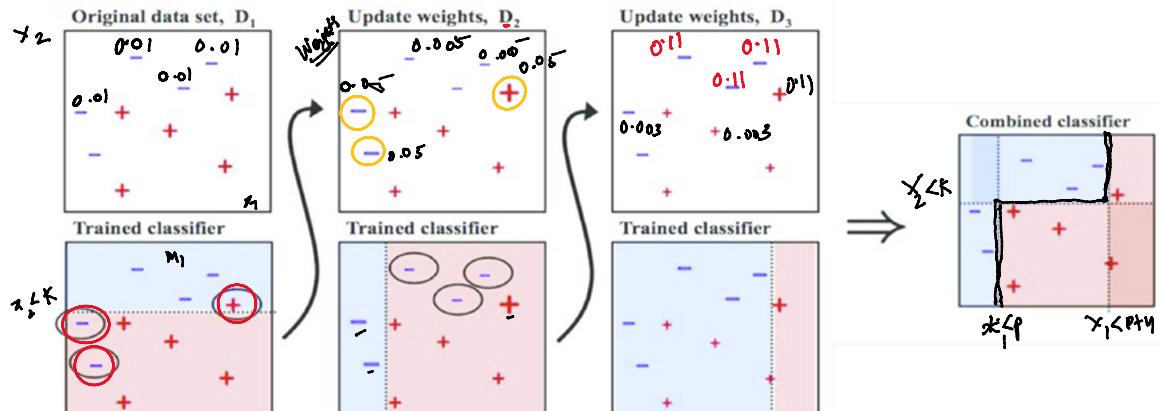
→ Sequential Process (my triple person marathon)

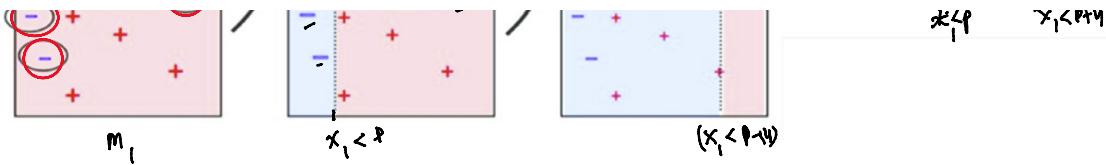


→ Adaptive Boosting (Ada boost)

→ Gradient Boosting (GBM), Xgboost, Extreme Gradient Boosting

Adaboost (decision stumps) — decision tree with only one split)
↑(weak learners)





'amount of say'

$$w_1 \propto \frac{1}{2} \log^9$$

$$w_2 \propto \frac{1}{2} \log \frac{1-0.1}{0.1}$$

Iteration 1 (Model 1) $M_1 = 0.9$

$N =$

100

$$w = \frac{1}{100} = 0.01$$

Initialize weights to all training points

$$w = \frac{1}{N}$$

Calculate error rate for each weak classifier

$$\epsilon = \sum_{\text{wrong}} w_i$$

$$\epsilon = 10 \times 0.01 = 0.1 \quad 0.9$$

(Model 2)

Pick Classifier with the lowest error rate

Now compute Voting Power for the classifier

$$\alpha = \frac{1}{2} \log \frac{1 - \epsilon}{\epsilon}$$

$$w_2 = \frac{1}{2} \log \frac{1 - 0.09}{0.09}$$

$$w_3 = \frac{1}{2} \log \frac{1 - \epsilon_3}{\epsilon_3}$$

Update weights of each point where the previous classifier went wrong

$$w_{\text{new}} = \begin{cases} \frac{w_{\text{old}}}{2(1-\epsilon)} & \text{if point classified correctly} \\ \frac{w_{\text{old}}}{2\epsilon} & \text{if point classified wrongly} \end{cases}$$

Append the classifier in the ensemble classifier and check if the classifier is good enough?

$$f(x_i) = \sum_{t=1}^T \alpha_t h_t(x_i)$$

$$w_1 = 1.09$$

$$w_2 = 0.81$$

$$w_3 = 0.76$$

$$M = \left[\frac{1.09 \times M_1 + 0.81 \times M_2}{1.09 + 0.81} \right]$$

$$M = \left[\frac{1.09 \times M_1 + 0.81 \times M_2 + 0.76 M_3}{1.09 + 0.81 + 0.76} \right]$$

$$w_{\text{new}} = \frac{w_{\text{old}}}{2(1-\epsilon)} \quad \text{if correctly classified}$$

$$w_{\text{new}} = \frac{w_{\text{old}}}{2\epsilon} \quad \text{if wrongly classified.}$$

$$w_{\text{new}} = \frac{0.01}{2(1-0.1)} = 0.005$$

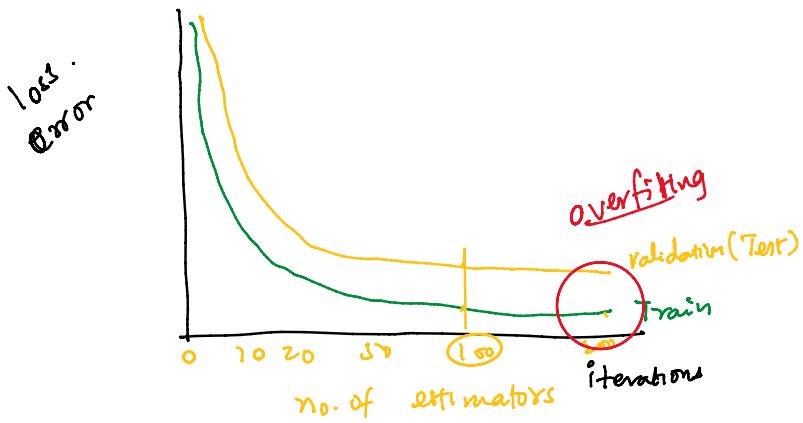
$$\frac{w_{\text{old}}}{2\epsilon} = \frac{0.01}{2 \times 0.1} = 0.05$$

$$w_{\text{new}} =$$

$$\text{Validation accuracy} = 0.91 - - -$$

$$w_{\text{new}} =$$

? N estimators :- No. of models



Classification: step by step

1. Initialize the weight $1/m$ for all $i=1..m$ examples
2. We create a decision stump for each variable and see how well each stump classifies samples to their target classes. We'd look at how many samples are correctly or incorrectly classified as Fit or Unfit for each individual stump.
3. We calculate the Gini index and build the tree
4. Total error for a stump is the sum of the weights associated with the incorrectly classified samples
5. We use the total error to determine the Amount of Say.
Amount of say = $\frac{1}{2} \log \frac{1-\epsilon_i}{\epsilon_i}$
6. We need to modify the weights by using the below formula
New Sample Weight = Sample Weight * $e^{\text{amount of say}}$
7. More weight is assigned to the incorrectly classified samples so that they're classified correctly in the next decision stump. Weight is also assigned to each classifier based on the accuracy of the classifier, which means high accuracy = high weight!
8. Reiterate from Step 2 until all the data points have been correctly classified, or the maximum iteration level has been reached

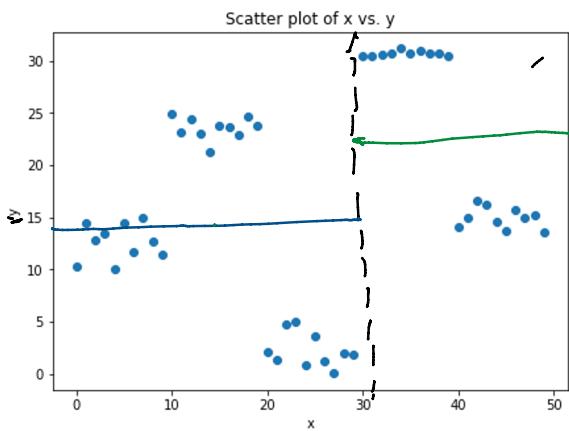
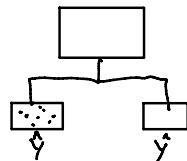
$$\eta = 0.5$$

$$\text{amount of say} = \frac{1}{2} \log \frac{1-\epsilon_i}{\epsilon_i}$$

$$\text{amount of say} = \eta \cdot \log \frac{1-\epsilon_i}{\epsilon_i}$$

η - Learning rate

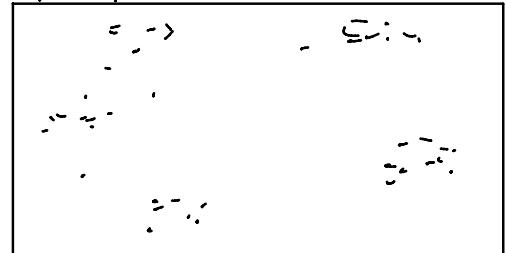
Gradient Boosting (Regression)



linear Regression

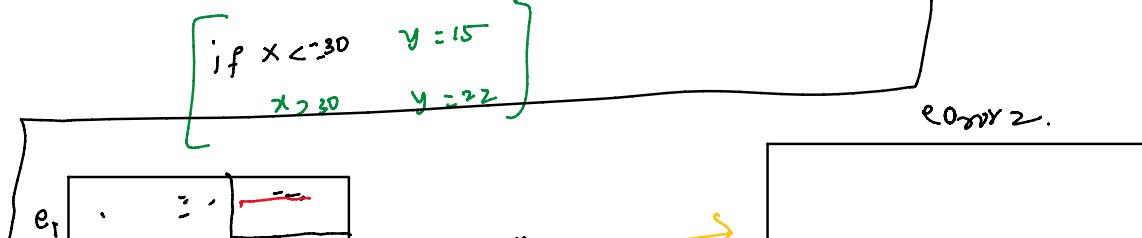
$$y = \beta_0 + \beta_1 x_1$$

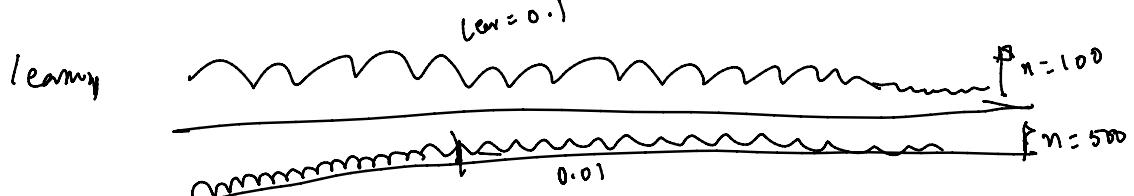
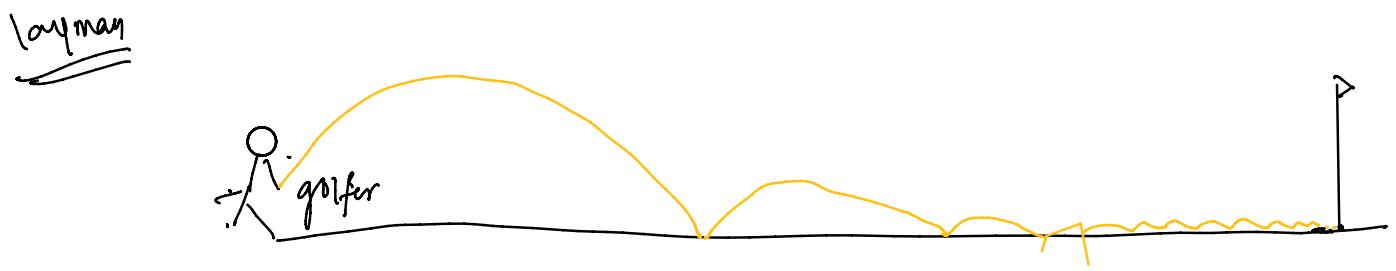
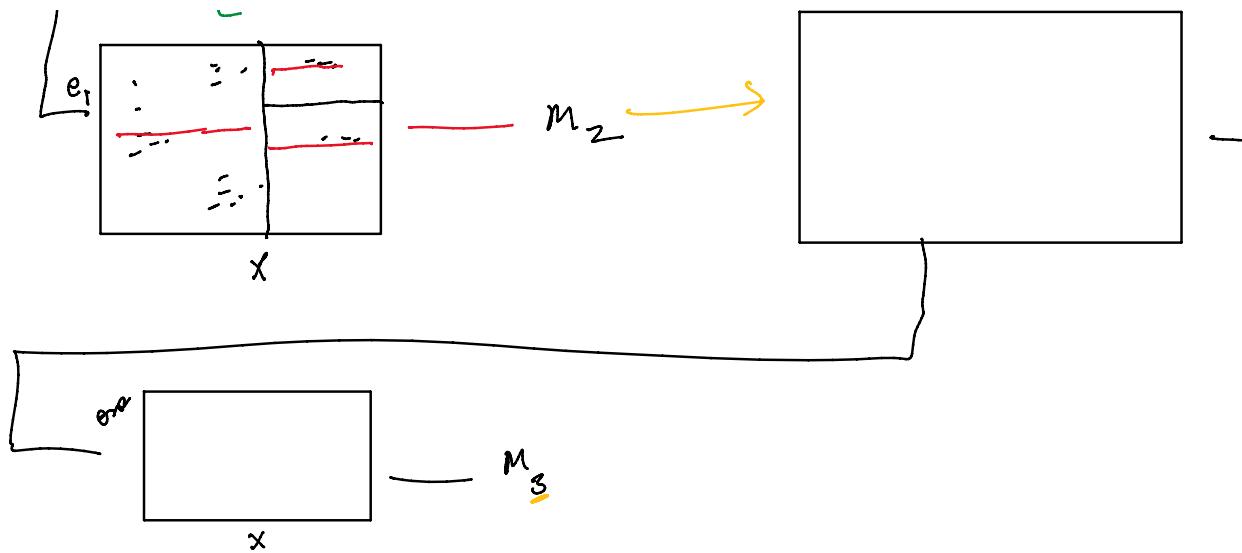
$$\beta = 0 \text{ errors}$$



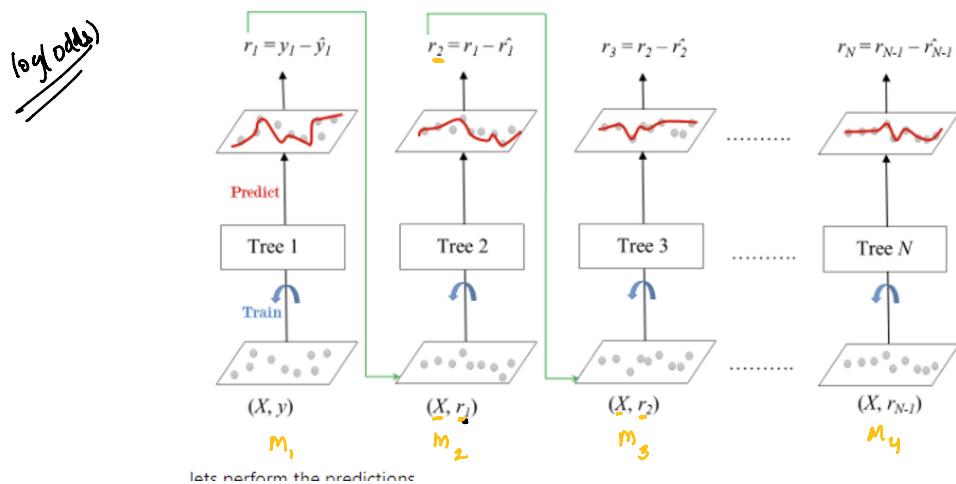
$$-M_1 \rightarrow e_{\text{err}}$$

$$\begin{cases} \text{if } x < 30 \\ x > 30 \end{cases} \quad \begin{cases} y = 15 \\ y = 22 \end{cases}$$





- where $r_i(w) = \text{previous base learner} \setminus r_m(w) = \text{current learner}$.



$$y = m_1 + m_2 + m_3 + m_4$$

Gradient descent

$$\text{loss} = L(y_i, \hat{y}_i)$$

$$\partial L = \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

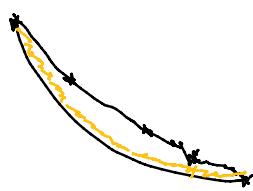
$$\hat{y}_i = y_i + \alpha \cdot \frac{\text{learning rate}}{\partial y}$$

$$= y_i + \alpha (-2)(y_i - \hat{y}_i)$$

$= y_i - \alpha \cdot 2 \cdot (y_i - \hat{y}_i)$

Gradient Boosting machine

$$\begin{aligned} n &= 10 \\ \eta &= 0.5 \\ \text{learn} &= 0.01 \\ \text{iter} &= 1000 \end{aligned}$$



disadvantage

Too slow

Bias variance tradeoff
High variance

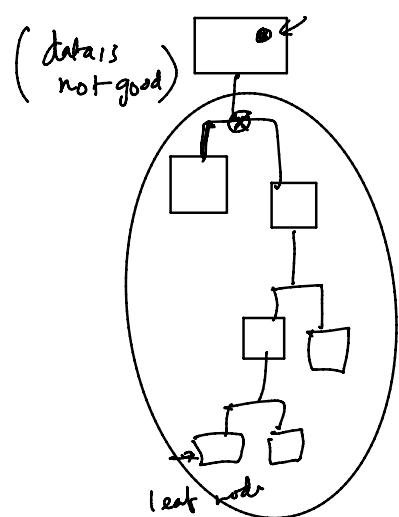
Xgboost - Whole bunch of boost functions

2. Parallelization (Tree - Entropy / Gini) ✓

3. Subsampling

4. Regularization

5. 'Gpu' execution ✓



Overfits - banana

Ensemble R.F., Boosting (Sub Sampling)

