

# Imbalanced\_Nov\_notebook

February 14, 2022

```
[2]: import pandas as pd
import numpy as np
from numpy import argmax

from datetime import date, time, timedelta
import pendulum

import matplotlib.pyplot as plt
import seaborn as sns

import category_encoders as ce
from sklearn.preprocessing import LabelEncoder, StandardScaler

from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.metrics import confusion_matrix, accuracy_score, precision_score, \
    recall_score, f1_score, precision_recall_curve
```

```
[90]: df = pd.read_csv("orders.csv")
```

```
[64]: df.head()
```

```
[64]: Unnamed: 0  order_id          created_at      city \
0            0   2166688  2019-10-01 00:21:37+00:00  SINGAPORE
1            1   2166694  2019-10-01 00:29:51+00:00  SINGAPORE
2            2   2166764  2019-10-01 00:55:24+00:00    DUBAI
3            3   2166991  2019-10-01 02:48:55+00:00  SINGAPORE
4            4   2167074  2019-10-01 03:27:07+00:00  SINGAPORE

          category_name  experience_date  product_id \
0      Singapore Zoo      2019-10-01          7360
1      Singapore Zoo      2019-10-01          7360
2          Dubai Frame      2019-12-07          8541
3  Singapore Cable Car      2019-10-01          7372
4  Universal Studios Singapore  2019-10-02          7442

          product_name  amount      device ... \
0  Singapore Zoo with Tram Ride  110.80  HIGH_END_MOBILE ...
1  Singapore Zoo with Tram Ride  110.80  HIGH_END_MOBILE ...
```

2	Dubai Frame Anytime Entry Tickets	32.67	HIGH_END_MOBILE	...
3	Singapore Cable Car Sky Pass: Round Trip	66.14	HIGH_END_MOBILE	...
4	Universal Studios Singapore Tickets	106.77	DESKTOP	...

	created_dow_sin	created_dow_cos	created_wom_sin	created_wom_cos	\
0	0.866025	0.5	0.951057	0.309017	
1	0.866025	0.5	0.951057	0.309017	
2	0.866025	0.5	0.951057	0.309017	
3	0.866025	0.5	0.951057	0.309017	
4	0.866025	0.5	0.951057	0.309017	

	experience_dom_sin	experience_dom_cos	experience_dow_sin	\
0	0.201299	0.979530	0.866025	
1	0.201299	0.979530	0.866025	
2	0.988468	0.151428	-0.866025	
3	0.201299	0.979530	0.866025	
4	0.394356	0.918958	0.866025	

	experience_dow_cos	experience_wom_sin	experience_wom_cos
0	0.5	0.951057	0.309017
1	0.5	0.951057	0.309017
2	0.5	0.587785	-0.809017
3	0.5	0.951057	0.309017
4	-0.5	0.951057	0.309017

[5 rows x 37 columns]

```
[91]: df = df.drop(["created_at", "experience_date", "Unnamed: 0"], axis = 1)
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 76829 entries, 0 to 76828
```

```
Data columns (total 34 columns):
```

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	order_id	76829 non-null	int64
1	city	76829 non-null	object
2	category_name	76743 non-null	object
3	product_id	76829 non-null	int64
4	product_name	76829 non-null	object
5	amount	76829 non-null	float64
6	device	76829 non-null	object
7	payment_id	76829 non-null	object
8	customer_ip	76829 non-null	object
9	customer_id	76829 non-null	object
10	payment_method	76829 non-null	object
11	payment_method_provider	76829 non-null	object
12	payment_method_bin	76778 non-null	float64

```

13 payment_method_type          76741 non-null object
14 payment_method_product       74731 non-null object
15 payment_method_card_category 74228 non-null object
16 payment_method_issuer_bank   73690 non-null object
17 payment_method_issuer_country 76730 non-null object
18 is_fraudulent                76829 non-null bool
19 time_diff                    76829 non-null float64
20 created_at_hod_sin           76829 non-null float64
21 created_at_hod_cos           76829 non-null float64
22 created_dom_sin              76829 non-null float64
23 created_dom_cos              76829 non-null float64
24 created_dow_sin              76829 non-null float64
25 created_dow_cos              76829 non-null float64
26 created_wom_sin              76829 non-null float64
27 created_wom_cos              76829 non-null float64
28 experience_dom_sin           76829 non-null float64
29 experience_dom_cos           76829 non-null float64
30 experience_dow_sin           76829 non-null float64
31 experience_dow_cos           76829 non-null float64
32 experience_wom_sin           76829 non-null float64
33 experience_wom_cos           76829 non-null float64
dtypes: bool(1), float64(17), int64(2), object(14)
memory usage: 19.4+ MB

```

```
[92]: df.columns
```

```

[92]: Index(['order_id', 'city', 'category_name', 'product_id', 'product_name',
          'amount', 'device', 'payment_id', 'customer_ip', 'customer_id',
          'payment_method', 'payment_method_provider', 'payment_method_bin',
          'payment_method_type', 'payment_method_product',
          'payment_method_card_category', 'payment_method_issuer_bank',
          'payment_method_issuer_country', 'is_fraudulent', 'time_diff',
          'created_at_hod_sin', 'created_at_hod_cos', 'created_dom_sin',
          'created_dom_cos', 'created_dow_sin', 'created_dow_cos',
          'created_wom_sin', 'created_wom_cos', 'experience_dom_sin',
          'experience_dom_cos', 'experience_dow_sin', 'experience_dow_cos',
          'experience_wom_sin', 'experience_wom_cos'],
          dtype='object')

```

```
[93]: df.is_fraudulent.value_counts(normalize=True)
```

```

[93]: False    0.986828
      True     0.013172
      Name: is_fraudulent, dtype: float64

```

```

[94]: df = df.drop(columns=["order_id", "payment_method", "payment_id",
    ↪ "product_name", "payment_method_bin"])
      df.head()

```

```

[94]:      city      category_name  product_id  amount  \
0  SINGAPORE      Singapore Zoo      7360  110.80
1  SINGAPORE      Singapore Zoo      7360  110.80
2    DUBAI      Dubai Frame      8541   32.67
3  SINGAPORE      Singapore Cable Car      7372   66.14
4  SINGAPORE  Universal Studios Singapore      7442  106.77

      device      customer_ip      customer_id  \
0  HIGH_END_MOBILE  vzH6NNH4iwxhe+eSr1L84w==  /8+13Ig7x5Jbiszc0f4AHw==
1  HIGH_END_MOBILE  vzH6NNH4iwxhe+eSr1L84w==  /8+13Ig7x5Jbiszc0f4AHw==
2  HIGH_END_MOBILE  cr7Phm0QLRXcaKa/bjpTDg==  gjP75WqjRZ9ZtSkHoFd+mw==
3  HIGH_END_MOBILE  ckoTvtWdM4zKRdFyDH4YcQ==  XDZdavE7L2SLdQ1zdp5FAw==
4    DESKTOP      5H/u0Mb8cMZxsXNDZ2QYw==  9M/VfPTvlko/2Zdtg5Pg1A==

      payment_method_provider  payment_method_type  payment_method_product  ...  \
0      Mastercard      Credit  World MasterCard® Card  ...
1      Mastercard      Credit  World MasterCard® Card  ...
2      Mastercard      Credit  Platinum MasterCard® Card  ...
3      Mastercard      Credit  Titanium MasterCard® Card  ...
4      Mastercard      Debit  Debit MasterCard® Card  ...

      created_dow_sin  created_dow_cos  created_wom_sin  created_wom_cos  \
0      0.866025      0.5      0.951057      0.309017
1      0.866025      0.5      0.951057      0.309017
2      0.866025      0.5      0.951057      0.309017
3      0.866025      0.5      0.951057      0.309017
4      0.866025      0.5      0.951057      0.309017

      experience_dom_sin  experience_dom_cos  experience_dow_sin  \
0      0.201299      0.979530      0.866025
1      0.201299      0.979530      0.866025
2      0.988468      0.151428     -0.866025
3      0.201299      0.979530      0.866025
4      0.394356      0.918958      0.866025

      experience_dow_cos  experience_wom_sin  experience_wom_cos
0      0.5      0.951057      0.309017
1      0.5      0.951057      0.309017
2      0.5      0.587785     -0.809017
3      0.5      0.951057      0.309017
4     -0.5      0.951057      0.309017

```

[5 rows x 29 columns]

```

[95]: if df.shape[0] == df.drop_duplicates().shape[0] :
      print('No duplicates Found')
      else:

```

```

duplicates = df.shape[0] - df.drop_duplicates().shape[0]
print('{} duplicates found'.format(duplicates))

```

6857 duplicates found

```
[96]: df.nunique()
```

```

[96]: city                2
      category_name       96
      product_id         353
      amount             8627
      device              5
      customer_ip       44208
      customer_id       51109
      payment_method_provider  5
      payment_method_type    8
      payment_method_product 144
      payment_method_card_category  2
      payment_method_issuer_bank 2046
      payment_method_issuer_country 151
      is_fraudulent        2
      time_diff           117
      created_at_hod_sin    24
      created_at_hod_cos    22
      created_dom_sin       31
      created_dom_cos       25
      created_dow_sin        7
      created_dow_cos        5
      created_wom_sin        6
      created_wom_cos        6
      experience_dom_sin     31
      experience_dom_cos     25
      experience_dow_sin      7
      experience_dow_cos      5
      experience_wom_sin      6
      experience_wom_cos      6
      dtype: int64

```

```

[97]: df = df.dropna(axis = 0, how= 'any',
                    subset = ['payment_method_issuer_country', 'payment_method_issuer_bank', 'payment_method_product',
                              'payment_method_card_category'])

```

```
[98]: df.is_fraudulent.value_counts(normalize=True)
```

```

[98]: False    0.985446
      True     0.014554
      Name: is_fraudulent, dtype: float64

```

```
[99]: df.isna().sum()
```

```
[99]: city                                0
      category_name                       75
      product_id                          0
      amount                              0
      device                              0
      customer_ip                         0
      customer_id                         0
      payment_method_provider             0
      payment_method_type                 0
      payment_method_product              0
      payment_method_card_category        0
      payment_method_issuer_bank          0
      payment_method_issuer_country       0
      is_fraudulent                      0
      time_diff                          0
      created_at_hod_sin                  0
      created_at_hod_cos                  0
      created_dom_sin                     0
      created_dom_cos                     0
      created_dow_sin                     0
      created_dow_cos                     0
      created_wom_sin                     0
      created_wom_cos                     0
      experience_dom_sin                  0
      experience_dom_cos                  0
      experience_dow_sin                  0
      experience_dow_cos                  0
      experience_wom_sin                  0
      experience_wom_cos                  0
      dtype: int64
```

```
[100]: list(df[df['category_name'].isna() == True]['product_id'].drop_duplicates())
```

```
[100]: [7364]
```

```
[101]: df['category_name'] = df['category_name'].fillna('other')
```

```
[102]: for column in list(df.columns) :
        if df[f"{column}"].dtypes == object :
            if len(df[f"{column}"].unique()) == len(df[f"{column}"].apply(lambda x :
→ x.lower()).unique()):
                print('no string duplicates in {}'.format(column))
            else:
                print('string duplicates in {}'.format(column))
```

```
no string duplicates in city
```

```

no string duplicates in category_name
no string duplicates in device
no string duplicates in customer_ip
no string duplicates in customer_id
no string duplicates in payment_method_provider
no string duplicates in payment_method_type
no string duplicates in payment_method_product
no string duplicates in payment_method_card_category
no string duplicates in payment_method_issuer_bank
no string duplicates in payment_method_issuer_country

```

```
[103]: X = df.drop(['is_fraudulent'], axis = 1)
      y = df['is_fraudulent']
```

```
[104]: le = LabelEncoder()
      X['city'] = le.fit_transform(X['city'])
      X['payment_method_card_category'] = le.
        ↳fit_transform(X['payment_method_card_category'])
      y = le.fit_transform(y)
```

```
[174]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
        ↳random_state = 0, stratify=y)
```

```
[175]: X_train.shape
```

```
[175]: (55628, 28)
```

```
[176]: X_test.shape
```

```
[176]: (13908, 28)
```

```
[177]: ce_target = ce.TargetEncoder(cols = ['customer_ip',
        ↳'customer_id', 'device', 'category_name', 'product_id',
        ↳
        ↳'payment_method_issuer_country', 'payment_method_issuer_bank', 'payment_method_provider',
        ↳
        ↳'payment_method_type', 'payment_method_product'])
      X_train = ce_target.fit_transform(X_train, y_train)
      X_test = ce_target.transform(X_test)
```

```
[178]: X_train.head()
```

```
[178]:
```

	city	category_name	product_id	amount	device	customer_ip	\
74012	0	0.000359	0.000363	4.08	0.001629	0.014561	
772	0	0.001008	0.000000	176.97	0.000000	0.003916	
36668	0	0.016409	0.000000	650.03	0.048906	0.000097	
29224	0	0.000000	0.000000	57.44	0.001629	0.014561	
15716	0	0.000000	0.000000	49.78	0.001629	0.003916	

	customer_id	payment_method_provider	payment_method_type	\
74012	0.014561	0.015633	0.005531	
772	0.003916	0.015633	0.005531	
36668	0.000036	0.012961	0.005531	
29224	0.014561	0.015633	0.019733	
15716	0.003916	0.015633	0.005531	

	payment_method_product	...	created_dow_sin	created_dow_cos	\
74012	5.523554e-04	...	-8.660254e-01	-0.5	
772	1.765252e-02	...	-8.660254e-01	0.5	
36668	0.000000e+00	...	8.660254e-01	-0.5	
29224	3.705241e-15	...	-2.449294e-16	1.0	
15716	5.523554e-04	...	1.224647e-16	-1.0	

	created_wom_sin	created_wom_cos	experience_dom_sin	\
74012	-9.510565e-01	0.309017	-0.790776	
772	5.877853e-01	-0.809017	0.651372	
36668	-2.449294e-16	1.000000	-0.848644	
29224	9.510565e-01	0.309017	0.937752	
15716	-5.877853e-01	-0.809017	0.485302	

	experience_dom_cos	experience_dow_sin	experience_dow_cos	\
74012	-0.612106	-8.660254e-01	-0.5	
772	-0.758758	-8.660254e-01	0.5	
36668	0.528964	1.224647e-16	-1.0	
29224	0.347305	-2.449294e-16	1.0	
15716	-0.874347	-8.660254e-01	-0.5	

	experience_wom_sin	experience_wom_cos
74012	-9.510565e-01	0.309017
772	5.877853e-01	-0.809017
36668	-2.449294e-16	1.000000
29224	9.510565e-01	0.309017
15716	-5.877853e-01	-0.809017

[5 rows x 28 columns]

```
[179]: X_train = X_train.values
       X_test = X_test.values
```

```
[180]: continuous_columns = ['amount', 'time_diff']
```

```
[181]: numerical_features = [df.columns.get_loc(c) for c in continuous_columns if c in_
    ↪ df]
```

```
[182]: numerical_features
```



```
[182]: [3, 14]
```

```
[183]: sc = StandardScaler() ## The issue was it is expecting an array
X_train[:,numerical_features] = sc.fit_transform(X_train[:,numerical_features])
X_test[:,numerical_features] = sc.transform(X_test[:,numerical_features])
```

```
[184]: print(X_train.shape)
print(y_train.shape)
```

```
(55628, 28)
```

```
(55628,)
```

```
[188]: from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import cross_validate

model = LogisticRegression(max_iter = 1000)# Restricting to max Iterations to
↳avoid convergence issues
scores = cross_validate(model, X_train, y_train, cv = 5, return_train_score=
↳True)
```

```
[189]: results = pd.DataFrame(scores)
```

```
[190]: results
```

```
[190]:   fit_time  score_time  test_score  train_score
0  0.622996    0.002003    0.993169    0.993281
1  0.518002    0.002001    0.993439    0.993416
2  0.568133    0.001000    0.994338    0.993101
3  0.610553    0.001000    0.992090    0.993618
4  0.559009    0.000993    0.993169    0.993439
```

```
[191]: results.test_score.mean()
```

```
[191]: 0.9932407919904183
```

```
[192]: results.train_score.mean()
```

```
[192]: 0.99337114261358
```

```
[193]: from sklearn.metrics import confusion_matrix, classification_report

model.fit(X_train, y_train)
conf_mat = confusion_matrix(y_train, model.predict(X_train))
print(conf_mat)
```

```
[[54683  135]
 [  234  576]]
```

```
[194]: tp = float(conf_mat[1][1])
tn = float(conf_mat[0][0])
fp = float(conf_mat[0][1])
fn = float(conf_mat[1][0])

print(classification_report(y_train, model.predict(X_train)))

spec = tn/(tn+fp)
sens = tp/(tp+fp)

g_mean = (spec*sens)**(0.5)

print(f"G-mean: {g_mean}")
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	54818
1	0.81	0.71	0.76	810
accuracy			0.99	55628
macro avg	0.90	0.85	0.88	55628
weighted avg	0.99	0.99	0.99	55628

G-mean: 0.8989613383285026

```
[195]: ## Test set performance
print(classification_report(y_test, model.predict(X_test)))
```

	precision	recall	f1-score	support
0	0.99	1.00	0.99	13706
1	0.69	0.43	0.53	202
accuracy			0.99	13908
macro avg	0.84	0.71	0.76	13908
weighted avg	0.99	0.99	0.99	13908

```
[135]: from imblearn.under_sampling import RandomUnderSampler

rus = RandomUnderSampler(random_state=0)
X_resampled, y_resampled = rus.fit_resample(X_train, y_train)
print(X_resampled.shape)
```

(1620, 28)

```
[137]: print(y_resampled[y_resampled == 1].shape, y_resampled[y_resampled == 0].shape)
```

(810,) (810,)

```
[138]: mets = ["f1", "precision", "recall", "accuracy"]
score_rus = cross_validate(model, X_resampled, y_resampled, cv = 5,
    ↪return_train_score = True, scoring = mets)
```

```
C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:814:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:814:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:814:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:814:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:814:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

regression

```
n_iter_i = _check_optimize_result(
```

```
[139]: pd.DataFrame(score_rus)
```

```
[139]:
```

	fit_time	score_time	test_f1	train_f1	test_precision	train_precision \
0	0.061522	0.003999	0.954407	0.962221	0.940120	0.961479
1	0.058538	0.004002	0.947368	0.967543	0.950311	0.969040
2	0.059002	0.004001	0.956522	0.968044	0.962500	0.977953
3	0.060680	0.004000	0.956522	0.964981	0.962500	0.973312
4	0.061591	0.004002	0.981132	0.963424	1.000000	0.971743

	test_recall	train_recall	test_accuracy	train_accuracy
0	0.969136	0.962963	0.953704	0.962191
1	0.944444	0.966049	0.947531	0.967593
2	0.950617	0.958333	0.956790	0.968364
3	0.950617	0.956790	0.956790	0.965278
4	0.962963	0.955247	0.981481	0.963735

```
[140]: model.fit(X_resampled, y_resampled)
prob = model.predict_proba(X_resampled)[: , 1]

f1_0 = []
f1_1 = []

thresh = [0.9, 0.8, 0.6, 0.5, 0.4, 0.3, 0.2, 0.1]
for i in thresh:

    preds = np.where(prob >= i, 1, 0)

    f1_0.append(f1_score(y_resampled, preds, pos_label = 0))
    f1_1.append(f1_score(y_resampled, preds))

plt.plot(thresh, f1_1)
plt.plot(thresh, f1_0)
plt.ylabel("F1-score")
plt.xlabel("thresh")
plt.legend(["Class1", "Class 0"])
plt.show()
```

C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:814:

ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.


Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```



```
[141]: test_prob = model.predict_proba(X_test)[: , 1]
preds = np.where(test_prob>0.4, 1, 0)
print(classification_report(y_test, preds))
```

	precision	recall	f1-score	support
0	1.00	0.95	0.98	13706
1	0.23	0.93	0.36	202
accuracy			0.95	13908
macro avg	0.61	0.94	0.67	13908
weighted avg	0.99	0.95	0.97	13908

```
[142]: from imblearn.under_sampling import TomekLinks
from collections import Counter

tom = TomekLinks()
```

```
X_tom, y_tom = tom.fit_resample(X_train, y_train)

print('Resampled dataset shape {}'.format(Counter(y_tom)))
```

Resampled dataset shape Counter({0: 54705, 1: 810})

```
[143]: score_tom = cross_validate(model, X_tom, y_tom, cv = 5, return_train_score =
↳ True, scoring = mets)
pd.DataFrame(score_tom)
```

C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:814:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:814:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:814:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:814:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```

regression
  n_iter_i = _check_optimize_result(
C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:814:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

regression

```

  n_iter_i = _check_optimize_result(

```

```

[143]:   fit_time  score_time  test_f1  train_f1  test_precision  train_precision  \
0  0.349674    0.010999  0.754839  0.769357         0.790541         0.825088
1  0.326775    0.009986  0.718644  0.765101         0.796992         0.838235
2  0.320925    0.010997  0.748344  0.764557         0.807143         0.843575
3  0.325865    0.010999  0.772414  0.746193         0.875000         0.825843
4  0.424207    0.009997  0.781145  0.747885         0.859259         0.827715

```

```

      test_recall  train_recall  test_accuracy  train_accuracy
0      0.722222      0.720679      0.993155      0.993695
1      0.654321      0.703704      0.992525      0.993695
2      0.697531      0.699074      0.993155      0.993718
3      0.691358      0.680556      0.994056      0.993245
4      0.716049      0.682099      0.994146      0.993290

```

```

[144]: model.fit(X_tom, y_tom)
print(classification_report(y_tom, model.predict(X_tom)))

```

```

              precision    recall  f1-score   support

         0           1.00        1.00        1.00        54705
         1           0.83        0.70        0.76         810

 accuracy                   0.99        55515
 macro avg              0.91        0.85        0.88        55515
 weighted avg           0.99        0.99        0.99        55515

```

```

C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:814:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

regression

```
n_iter_i = _check_optimize_result(
```

```
[145]: print(classification_report(y_test, model.predict(X_test)))
```

	precision	recall	f1-score	support
0	0.99	1.00	0.99	13706
1	0.71	0.48	0.57	202
accuracy			0.99	13908
macro avg	0.85	0.74	0.78	13908
weighted avg	0.99	0.99	0.99	13908

```
[146]: from imblearn.over_sampling import SMOTE
```

```
smt = SMOTE()  
X_sm, y_sm = smt.fit_resample(X_train, y_train)  
  
print('Resampled dataset shape {}'.format(Counter(y_sm)))
```

```
Resampled dataset shape Counter({0: 54818, 1: 54818})
```

```
[147]: score_sm = cross_validate(model, X_sm, y_sm, cv = 5, return_train_score = True,  
    ↳ scoring = mets)  
pd.DataFrame(score_sm)
```

```
C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:814:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:814:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:814:  
ConvergenceWarning: lbfgs failed to converge (status=1):
```



STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:814:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:814:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
[147]:   fit_time  score_time  test_f1  train_f1  test_precision  train_precision  \
0  0.656133    0.017996  0.963469  0.965233         0.968396         0.968848
1  0.773557    0.020001  0.973717  0.974952         0.979332         0.979405
2  0.760311    0.017999  0.969825  0.970276         0.976688         0.975566
3  0.767751    0.019963  0.972310  0.971327         0.972620         0.972437
4  0.927605    0.020999  0.975299  0.973855         0.974677         0.974344
```

```
   test_recall  train_recall  test_accuracy  train_accuracy
0    0.958592    0.961645    0.963654    0.965362
1    0.968166    0.970539    0.973868    0.975065
2    0.963058    0.965044    0.970037    0.970436
3    0.971999    0.970219    0.972317    0.971360
4    0.975921    0.973366    0.975282    0.973868
```

```
[148]: model.fit(X_sm, y_sm)
```

```
C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:814:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
[148]: LogisticRegression()
```

```
[149]: print(classification_report(y_test, model.predict(X_test)))
```

	precision	recall	f1-score	support
0	1.00	0.97	0.99	13706
1	0.33	0.90	0.48	202
accuracy			0.97	13908
macro avg	0.66	0.93	0.73	13908
weighted avg	0.99	0.97	0.98	13908

```
[150]: from imblearn.over_sampling import ADASYN
```

```
ada = ADASYN(sampling_strategy=0.3)
```

```
X_ada, y_ada = ada.fit_resample(X_train, y_train)
```

```
print('Resampled dataset shape {}'.format(Counter(y_ada)))
```

Resampled dataset shape Counter({0: 54818, 1: 16520})

```
[151]: score_ada = cross_validate(model, X_ada, y_ada, cv = 5, return_train_score =  
↳ True, scoring = mets)  
pd.DataFrame(score_ada)
```

C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:814:

ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:814:

ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  

```
n_iter_i = _check_optimize_result(
```

C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:814:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  

```
n_iter_i = _check_optimize_result(
```

C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:814:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  

```
n_iter_i = _check_optimize_result(
```

C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:814:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  

```
n_iter_i = _check_optimize_result(
```

```
[151]:
```

	fit_time	score_time	test_f1	train_f1	test_precision	train_precision	\
0	0.451033	0.013000	0.918617	0.924161	0.941532	0.946175	
1	0.405699	0.011984	0.951282	0.929850	0.954033	0.954796	
2	0.459005	0.016001	0.898076	0.926400	0.946064	0.946475	
3	0.428010	0.013991	0.930161	0.921302	0.943631	0.936943	
4	0.421019	0.012982	0.929262	0.919939	0.942421	0.945238	
	test_recall	train_recall	test_accuracy	train_accuracy			
0	0.896792	0.903148	0.963204	0.965674			
1	0.948547	0.906174	0.977502	0.968337			

2	0.854722	0.907158	0.955074	0.966620
3	0.917070	0.906174	0.968108	0.964150
4	0.916465	0.895959	0.967688	0.963887

```
[152]: model.fit(X_ada, y_ada)
```

```
C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:814:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
[152]: LogisticRegression()
```

```
[153]: prob = model.predict_proba(X_ada)[: , 1]
```

```
f1_0 = []
```

```
f1_1 = []
```

```
thresh = [0.9, 0.8, 0.6, 0.5, 0.4, 0.3, 0.2, 0.1]
```

```
for i in thresh:
```

```
    preds = np.where(prob >= i, 1, 0)
```

```
    f1_0.append(f1_score(y_ada, preds, pos_label = 0))
```

```
    f1_1.append(f1_score(y_ada, preds))
```

```
plt.plot(thresh, f1_1)
```


```
plt.plot(thresh, f1_0)
```

```
plt.ylabel("F1-score")
```

```
plt.xlabel("thresh")
```

```
plt.legend(["Class1", "Class 0"])
```

```
plt.show()
```



```
[154]: test_prob = model.predict_proba(X_test)[: , 1]
preds = np.where(test_prob>0.4, 1, 0)
print(classification_report(y_test, preds))
```

	precision	recall	f1-score	support
0	1.00	0.98	0.99	13706
1	0.44	0.90	0.59	202
accuracy			0.98	13908
macro avg	0.72	0.94	0.79	13908
weighted avg	0.99	0.98	0.98	13908

```
[155]: from imblearn.combine import SMOTEENN

smenn = SMOTEENN()
X_smenn, y_smenn = smenn.fit_resample(X_train, y_train)

print('Resampled dataset shape {}'.format(Counter(y_smenn)))
```

Resampled dataset shape Counter({1: 54055, 0: 52925})

```
[156]: score_smenn = cross_validate(model, X_smenn, y_smenn, cv = 5,
↳return_train_score = True, scoring = mets)
```

```
pd.DataFrame(score_smenn)
```

```
C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:814:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:814:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:814:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:814:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:814:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
[156]: fit_time score_time test_f1 train_f1 test_precision train_precision \
0 0.648059 0.016999 0.986049 0.986201 0.991490 0.989889
1 0.720929 0.020001 0.974967 0.976153 0.984624 0.987529
2 0.760998 0.017999 0.969358 0.970148 0.977611 0.976569
3 0.767568 0.018998 0.986978 0.986581 0.992517 0.990536
4 0.671150 0.018998 0.976783 0.976145 0.978733 0.980025

test_recall train_recall test_accuracy train_accuracy
0 0.980668 0.982541 0.985979 0.986107
1 0.965498 0.965036 0.974949 0.976175
2 0.961243 0.963810 0.969293 0.970029
3 0.981500 0.982657 0.986913 0.986493
4 0.974840 0.972297 0.976584 0.975989
```

```
[157]: model.fit(X_smenn, y_smenn)
```

C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:814:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
[157]: LogisticRegression()
```

```
[158]: print(classification_report(y_test, model.predict(X_test)))
```

	precision	recall	f1-score	support
0	1.00	0.97	0.98	13706
1	0.29	0.92	0.44	202
accuracy			0.97	13908
macro avg	0.65	0.94	0.71	13908
weighted avg	0.99	0.97	0.97	13908

```
[159]: from imblearn.combine import SMOTETomek
```

```
smtk = SMOTETomek()
```

```
X_smtk, y_smtk = smtk.fit_resample(X_train, y_train)

print('Resampled dataset shape {}'.format(Counter(y_smtk)))
```

Resampled dataset shape Counter({0: 54766, 1: 54766})

```
[160]: score_smtk = cross_validate(model, X_smtk, y_smtk, cv = 5, return_train_score =
↳ True, scoring = mets)
pd.DataFrame(score_smtk)
```

C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:814:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:814:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:814:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:814:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)



```

regression
  n_iter_i = _check_optimize_result(
C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:814:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```

regression
  n_iter_i = _check_optimize_result(

```

```

[160]:   fit_time  score_time  test_f1  train_f1  test_precision  train_precision \
0  0.646359    0.016999  0.974710  0.974687         0.981843         0.981419
1  0.650550    0.017052  0.972908  0.974079         0.978751         0.979417
2  0.673408    0.019144  0.975015  0.975374         0.979190         0.979426
3  0.752331    0.019999  0.965362  0.964801         0.974156         0.972850
4  0.801777    0.018467  0.973707  0.972210         0.975357         0.975483

```

```

      test_recall  train_recall  test_accuracy  train_accuracy
0      0.967680      0.968046      0.974894      0.974859
1      0.967135      0.968799      0.973068      0.974220
2      0.970876      0.971356      0.975121      0.975475
3      0.956724      0.956885      0.965672      0.965090
4      0.972062      0.968959      0.973751      0.972303

```

```

[161]: model.fit(X_smtk, y_smtk)

```

```

C:\Users\sci\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:814:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```

regression
  n_iter_i = _check_optimize_result(

```

```

[161]: LogisticRegression()

```

```

[162]: print(classification_report(y_test, model.predict(X_test)))

```

```

              precision    recall  f1-score   support

0               1.00        0.98        0.99       13706
1               0.36        0.91        0.52         202

```

accuracy			0.98	13908
macro avg	0.68	0.94	0.75	13908
weighted avg	0.99	0.98	0.98	13908

[ ]: